

Nicholas Bindela

New York, New York | (914) 539-1544 | njb2163@columbia.edu | <https://njb2163.github.io/>

EDUCATION

Columbia University

M.S. in Computer Science, Concentration in Machine Learning GPA: 4.1/4.0

- Relevant coursework: Databases, Algorithms, AI, Applied ML, Computer Vision, Generative LLMs

New York, NY
Dec 2025

Bucknell University

B.S. in Computer Science and Engineering, Minor in Economics GPA: 3.7/4.0

- Relevant coursework: Linear Algebra, Discrete Structures, Probability & Statistics

Lewisburg, PA
May 2021

LANGUAGE AND IT SKILLS

Languages & Frameworks: Python, SQL, Next.js, Node.js, React.js, Javascript, TailwindCSS, HTML5, CSS3

Python Libraries: Flask, FastAPI, NumPy, Pandas, Matplotlib, Scikit-learn, TensorFlow, PyTorch

Databases: PostgreSQL, MySQL, MongoDB

Tools & Platforms: AWS, Terraform, Jenkins, Git, Docker

AWS Services: EMR, EC2, S3, Route 53, DynamoDB, Lambda, ECS, SQS, SNS, Kinesis, CloudWatch, EventBridge, AuroraDB

WORK EXPERIENCE

Columbia University and Stealth Startup

Columbia Build Lab Course, Technical Lead

- Partnered with the CEO to translate the business vision into a functional MVP, defining scope, technical requirements, and product roadmap. Managed a 3-engineer team by creating Epics/User Stories, distributing tasks, and providing technical guidance
- Architected and built the full-stack application: a Next.js frontend with Tailwind CSS, a Python backend, and Firebase resources. Set up GitHub version control, CI/CD pipeline with Google Cloud Deploy + Firebase Hosting, local integration testing, and REST APIs for frontend integration

New York, NY
Jan 2025 - Present

Chewy

Software Engineer II

- Spearheaded the design and implementation of new services on the supply chain platform, leading architecture, development, testing, and adoption of client-facing features while ensuring smooth integration with existing systems
- Built a multi-service architecture on AWS ECS, connecting a frontend to a backend with Aurora Postgres, orchestrating ML workloads on EMR, and enabling event-driven processing with DynamoDB and Kinesis for real-time status updates and communication

Boston, MA
Oct 2024 - Sep 2025

Software Engineer I

- Built a service to load input data to S3, utilized Crawlers to make data available to query in Athena, created EventBridge schedules to trigger 8+ ML models to run on EMR
- Implemented Lambdas to process output data from ML models and then send to external systems. Data is used to send Purchase Orders to over 100 vendors each day

Sep 2022 - Oct 2024

Columbia University and Ask2AI

AI in Finance Course, Machine Learning Research Assistant

- Led development and evaluation of predictive models for commercial loan defaults, training and testing over 10 linear and tree-based models to improve prediction accuracy
- Improved previous best AUC by 11%, achieving a new high of 0.96—surpassing prior benchmarks through advanced feature engineering, data preprocessing, and hyperparameter tuning
- Utilized Shapley Adaptive Reasoning to enhance model interpretability, revealing a commercial loanee's industry—particularly in Technology—significantly impacts default probability

New York, NY
Jun 2024 - Aug 2024

FAST Enterprises

Implementation Consultant

- Implemented a tax refund subsystem using C# and MySQL, resolved refund-related issues during two major software rollouts, and ran weekly client meetings to refine software functionality based on feedback

Hartford, CT
Sep 2021 - Aug 2022

Bucknell University

AI and Cognitive Science Research Assistant

- Conducted NSF-funded research over 12 months to develop a cognitive agent API using Python and Common-LISP, enabling human-like agents to learn from simulated environments and perform simple tasks

Lewisburg, PA
Jun 2019 - May 2020

TECHNICAL PROJECTS

End-to-End Social Media App: UI/UX & Fullstack Development

- Designed and developed a social media application by consulting stakeholders, prototyping 30 potential solutions, creating low and high-fidelity screens, and implementing a React front end with a Flask back end

Sep 2024 - Dec 2024

Predicting Spotify Track Popularity Using Regression Models

- Built and optimized regression models to analyze a dataset of 30,000 Spotify tracks, achieving a 9% improvement in MAE compared to baseline models and identified key predictors danceability, energy, and valence, with tree based models having best performance

Sep 2024 - Dec 2024